

Splitting the dataset

Choonghyun Ryu

2020-01-23

Preface

To develop a classification model, the original data must be divided into train data set and test data set. You should do the following:

- Cleansing the dataset
- **Split the data into a train set and a test set**
 - Split the `data.frame` or `tbl_df` into a train set and a test set
 - Compare dataset
 - * Comparison of categorical variables
 - * Comparison of numeric variables
 - * Diagnosis of train set and test set
 - Extract train/test dataset
 - * Extract train set or test set
 - * Extract the data to fit the model
- Modeling and Evaluate, Predict

The `alookr` package makes these steps fast and easy:

Data: Credit Card Default Data

Default of ISLR package is a simulated data set containing information on ten thousand customers. The aim here is to predict which customers will default on their credit card debt.

A data frame with 10000 observations on the following 4 variables.:

- `default` : factor. A factor with levels No and Yes indicating whether the customer defaulted on their debt
- `student`: factor. A factor with levels No and Yes indicating whether the customer is a student
- `balance`: numeric. The average balance that the customer has remaining on their credit card after making their monthly payment
- `income` : numeric. Income of customer

```
# Credit Card Default Data
head(ISLR::Default)
  default student  balance  income
1      No      No  729.5265 44361.625
2      No     Yes  817.1804 12106.135
3      No      No 1073.5492 31767.139
4      No      No  529.2506 35704.494
5      No      No  785.6559 38463.496
6      No     Yes  919.5885  7491.559

# structure of dataset
str(ISLR::Default)
'data.frame':  10000 obs. of  4 variables:
 $ default: Factor w/ 2 levels "No","Yes": 1 1 1 1 1 1 1 1 1 1 ...
 $ student: Factor w/ 2 levels "No","Yes": 1 2 1 1 1 2 1 2 1 1 ...
 $ balance: num  730 817 1074 529 786 ...
 $ income : num  44362 12106 31767 35704 38463 ...
```

```
# summary of dataset
summary(ISLR::Default)
default      student      balance      income
No :9667      No :7056      Min.   : 0.0      Min.   : 772
Yes: 333      Yes:2944      1st Qu.: 481.7    1st Qu.:21340
                        Median : 823.6    Median :34553
                        Mean   : 835.4    Mean   :33517
                        3rd Qu.:1166.3    3rd Qu.:43808
                        Max.   :2654.3    Max.   :73554
```

Split dataset

`split_by()` splits the `data.frame` or `tbl_df` into a training set and a test set.

Split dataset with `split_by()`

The `split_df` class is created, which contains the split information and criteria to separate the training and the test set.

```
library(alookr)
library(dplyr)

# Generate data for the example
sb <- ISLR::Default %>%
  split_by(default, seed = 6534)

sb
# A tibble: 10,000 x 5
# Groups:   split_flag [2]
  default student balance income split_flag
  <fct>    <fct>    <dbl>  <dbl> <chr>
1 No      No      730.  44362. train
2 No      Yes     817.  12106. train
3 No      No     1074.  31767. train
4 No      No     529.  35704. train
5 No      No     786.  38463. test
6 No      Yes     920.   7492. train
# ... with 9,994 more rows
```

The attributes of the `split_df` class are as follows.:

- `split_seed` : integer. random seed used for splitting
- `target` : character. the name of the target variable
- `binary` : logical. whether the target variable is binary class
- `minority` : character. the name of the minority class
- `majority` : character. the name of the majority class
- `minority_rate` : numeric. the rate of the minority class
- `majority_rate` : numeric. the rate of the majority class

```
sb_attr <- attributes(sb)

# The third attribute, row.names, is a vector that is very long and excluded from the output.
sb_attr[-3]
$names
[1] "default"      "student"      "balance"      "income"      "split_flag"
```

```

$class
[1] "split_df" "grouped_df" "tbl_df" "tbl" "data.frame"

$groups
# A tibble: 2 x 2
  split_flag .rows
  <chr>      <list>
1 test      <int [3,000]>
2 train     <int [7,000]>

$split_seed
[1] 6534

$target
default
"default"

$binary
[1] TRUE

$minority
[1] "Yes"

$majority
[1] "No"

$minority_rate
Yes
0.0333

$majority_rate
No
0.9667

```

`summary()` summarizes the information of two datasets splitted by `split_by()`.

```

summary(sb)
** Split train/test set information **
+ random seed      : 6534
+ split data
  - train set count : 7000
  - test set count  : 3000
+ target variable  : default
  - minority class  : Yes (0.033300)
  - majority class  : No (0.966700)

```

Compare dataset

Train data and test data should be similar. If the two datasets are not similar, the performance of the predictive model may be reduced.

`alookr` provides a function to compare the similarity between train dataset and test dataset.

If the two data sets are not similar, the train dataset and test dataset should be splitted again from the original data.

Comparison of categorical variables with `compare_category()`

Compare the statistics of the categorical variables of the train set and test set included in the “split_df” class.

```
sb %>%
  compare_category()
# A tibble: 4 x 5
  variable level train test abs_diff
  <chr>    <fct> <dbl> <dbl>    <dbl>
1 default No    96.7  96.7    0.00476
2 default Yes    3.33  3.33    0.00476
3 student No    70.0  71.8    1.77
4 student Yes    30.0  28.2    1.77

# compare variables that are character data types.
sb %>%
  compare_category(add_character = TRUE)
# A tibble: 4 x 5
  variable level train test abs_diff
  <chr>    <fct> <dbl> <dbl>    <dbl>
1 default No    96.7  96.7    0.00476
2 default Yes    3.33  3.33    0.00476
3 student No    70.0  71.8    1.77
4 student Yes    30.0  28.2    1.77

# display marginal
sb %>%
  compare_category(margin = TRUE)
# A tibble: 6 x 5
  variable level train test abs_diff
  <chr>    <fct> <dbl> <dbl>    <dbl>
1 default No    96.7  96.7    0.00476
2 default Yes    3.33  3.33    0.00476
3 default <Total> 100    100    0.00952
4 student No    70.0  71.8    1.77
5 student Yes    30.0  28.2    1.77
6 student <Total> 100    100    3.54

# student variable only
sb %>%
  compare_category(student)
# A tibble: 2 x 5
  variable level train test abs_diff
  <chr>    <fct> <dbl> <dbl>    <dbl>
1 student No    70.0  71.8    1.77
2 student Yes    30.0  28.2    1.77

sb %>%
  compare_category(student, margin = TRUE)
# A tibble: 3 x 5
  variable level train test abs_diff
  <chr>    <fct> <dbl> <dbl>    <dbl>
1 student No    70.0  71.8    1.77
2 student Yes    30.0  28.2    1.77
3 student <Total> 100    100    3.54
```

`compare_category()` returns `tbl_df`, where the variables have the following.:

- `variable` : character. categorical variable name
- `level` : factor. level of categorical variables
- `train` : numeric. the relative frequency of the level in the train set
- `test` : numeric. the relative frequency of the level in the test set
- `abs_diff` : numeric. the absolute value of the difference between two relative frequencies

Comparison of numeric variables with `compare_numeric()`

Compare the statistics of the numerical variables of the train set and test set included in the “split_df” class.

```
sb %>%
  compare_numeric()
# A tibble: 2 x 7
  variable train_mean test_mean train_sd test_sd train_z test_z
  <chr>      <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
1 balance      836.      834.     487.     477.     1.72     1.75
2 income    33446.    33684.   13437.   13101.    2.49     2.57

# balance variable only
sb %>%
  compare_numeric(balance)
# A tibble: 1 x 7
  variable train_mean test_mean train_sd test_sd train_z test_z
  <chr>      <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
1 balance      836.      834.     487.     477.     1.72     1.75
```

`compare_numeric()` returns `tbl_df`, where the variables have the following.:

- `variable` : character. numeric variable name
- `train_mean` : numeric. arithmetic mean of train set
- `test_mean` : numeric. arithmetic mean of test set
- `train_sd` : numeric. standard deviation of train set
- `test_sd` : numeric. standard deviation of test set
- `train_z` : numeric. the arithmetic mean of the train set divided by the standard deviation
- `test_z` : numeric. the arithmetic mean of the test set divided by the standard deviation


Comparison plot with `compare_plot()`

Plot compare information of the train set and test set included in the “split_df” class.


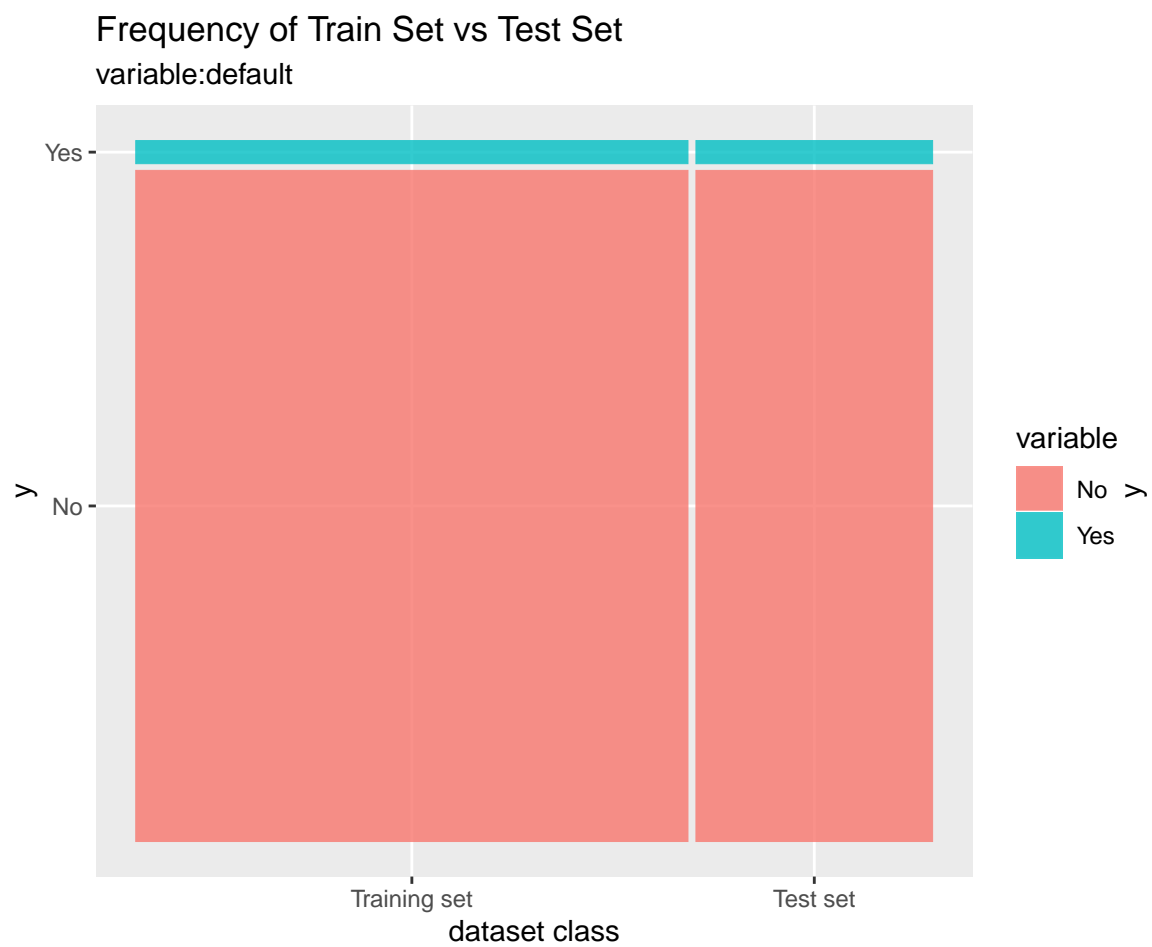
```
# income variable only
sb %>%
  compare_plot("income")
```

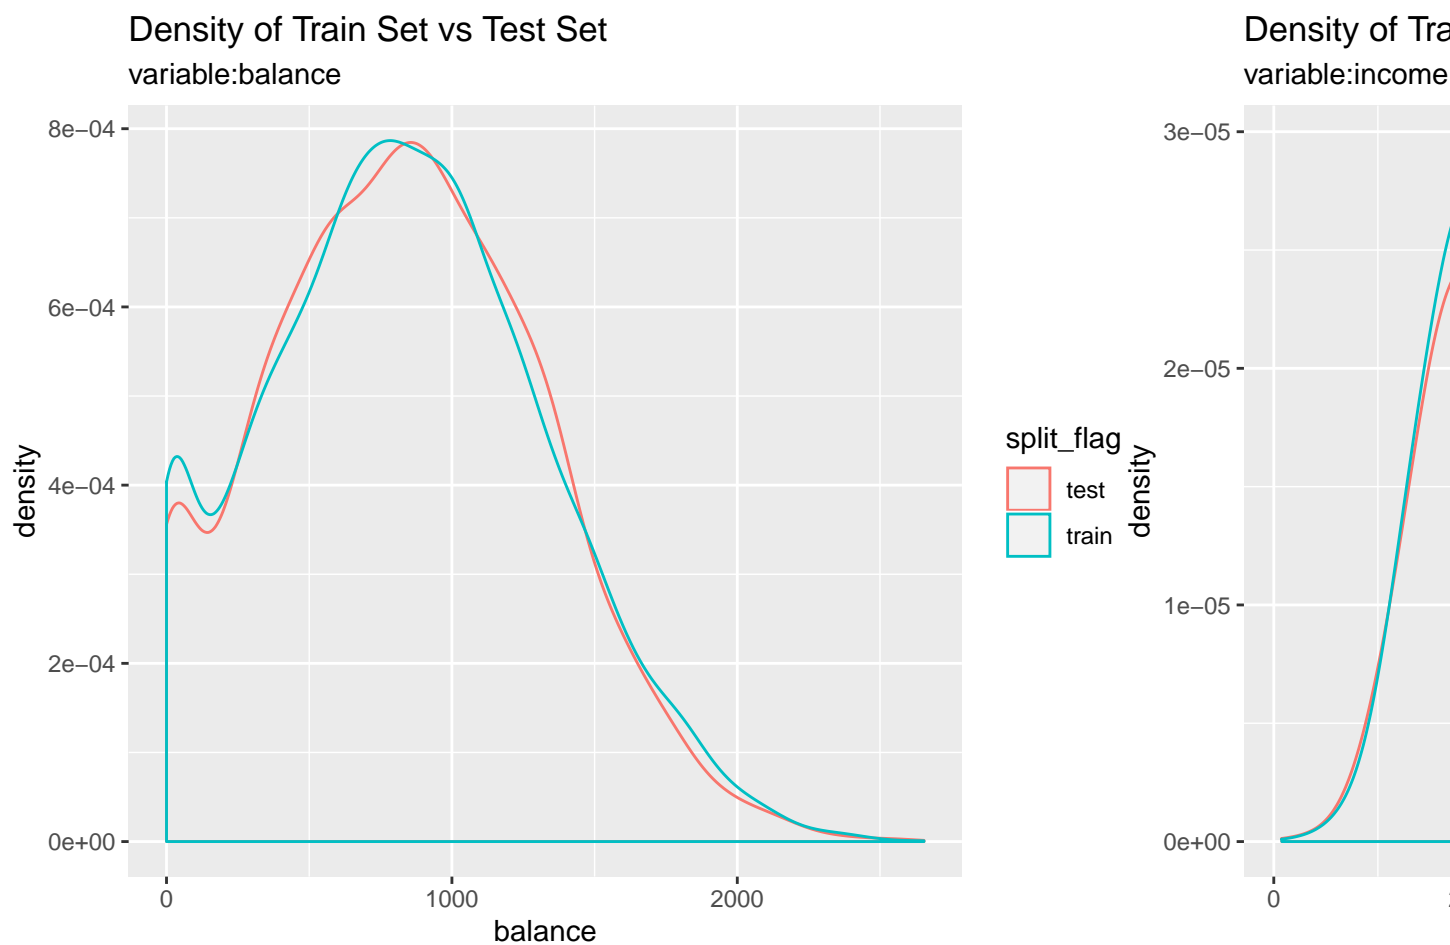
Density of Train Set vs Test Set

variable:income



```
# all variables  
sb %>%  
  compare_plot()
```





Diagnosis of train set and test set with `compare_diag()`

Diagnosis of similarity between datasets splitted by train set and set included in the “split_df” class.

```
defaults <- ISLR::Default
defaults$id <- seq(NROW(defaults))

set.seed(1)
defaults[sample(seq(NROW(defaults)), 3), "student"] <- NA
set.seed(2)
defaults[sample(seq(NROW(defaults)), 10), "balance"] <- NA

sb_2 <- defaults %>%
  split_by(default)

sb_2 %>%
  compare_diag()
* Detected diagnose missing value
- student
- balance

* Detected diagnose missing levels
- student
$missing_value
```



```

# A tibble: 2 x 4
  variables train_misscount train_missrate test_missrate
  <chr>          <int>          <dbl>          <dbl>
1 student              3          0.0429          NA
2 balance              5          0.0714          0.167

$single_value
# A tibble: 0 x 3
# ... with 3 variables: variables <chr>, train_uniq <lgl>, test_uniq <lgl>

$uniq_rate
# A tibble: 0 x 5
# ... with 5 variables: variables <chr>, train_uniqcount <int>,
#   train_uniqrate <dbl>, test_uniqcount <int>, test_uniqrate <dbl>

$missing_level
# A tibble: 1 x 4
  variables n_levels train_missing_nlevel test_missing_nlevel
  <chr>          <int>          <int>          <int>
1 student              3              0              1

sb_2 %>%
  compare_diag(add_character = TRUE)
* Detected diagnose missing value
- student
- balance

* Detected diagnose missing levels
- student

$missing_value
# A tibble: 2 x 4
  variables train_misscount train_missrate test_missrate
  <chr>          <int>          <dbl>          <dbl>
1 student              3          0.0429          NA
2 balance              5          0.0714          0.167

$single_value
# A tibble: 0 x 3
# ... with 3 variables: variables <chr>, train_uniq <lgl>, test_uniq <lgl>

$uniq_rate
# A tibble: 0 x 5
# ... with 5 variables: variables <chr>, train_uniqcount <int>,
#   train_uniqrate <dbl>, test_uniqcount <int>, test_uniqrate <dbl>

$missing_level
# A tibble: 1 x 4
  variables n_levels train_missing_nlevel test_missing_nlevel
  <chr>          <int>          <int>          <int>
1 student              3              0              1

sb_2 %>%
  compare_diag(uniq_thres = 0.0005)

```

```

* Detected diagnose missing value
- student
- balance

* Detected diagnose many unique value
- default
- student

* Detected diagnose missing levels
- student
$missing_value
# A tibble: 2 x 4
  variables train_misscount train_missrate test_missrate
  <chr>          <int>          <dbl>          <dbl>
1 student              3          0.0429          NA
2 balance              5          0.0714          0.167

$single_value
# A tibble: 0 x 3
# ... with 3 variables: variables <chr>, train_uniq <lgl>, test_uniq <lgl>

$uniq_rate
# A tibble: 2 x 5
  variables train_uniqcount train_uniqrate test_uniqcount test_uniqrate
  <chr>          <int>          <dbl>          <int>          <dbl>
1 default              NA              NA              2          0.000667
2 student              NA              NA              2          0.000667

$missing_level
# A tibble: 1 x 4
  variables n_levels train_missing_nlevel test_missing_nlevel
  <chr>          <int>          <int>          <int>
1 student              3              0              1

```

Extract train/test dataset

If you compare the train set with the test set and find that the two datasets are similar, extract the data from the `split_df` object.

Extract train set or test set with `extract_set()`

Extract train set or test set from `split_df` class object.

```

train <- sb %>%
  extract_set(set = "train")

test <- sb %>%
  extract_set(set = "test")

dim(train)
[1] 7000  4

dim(test)
[1] 3000  4

```

Extract the data to fit the model with `sampling_target()`

In a target class, the ratio of the majority class to the minority class is not similar and the ratio of the minority class is very small, which is called the **imbalanced class**.

If target variable is an imbalanced class, the characteristics of the majority class are actively reflected in the model. This model implies an error in predicting the minority class as the majority class. So we have to make the train dataset a balanced class.

`sampling_target()` performs sampling on the train set of `split_df` to resolve the imbalanced class.

```
# under-sampling with random seed
under <- sb %>%
  sampling_target(seed = 1234L)

under %>%
  count(default)
# A tibble: 2 x 2
  default      n
  <fct>    <int>
1 No       233
2 Yes      233

# under-sampling with random seed, and minority class frequency is 40%
under40 <- sb %>%
  sampling_target(seed = 1234L, perc = 40)

under40 %>%
  count(default)
# A tibble: 2 x 2
  default      n
  <fct>    <int>
1 No       349
2 Yes      233

# over-sampling with random seed
over <- sb %>%
  sampling_target(method = "ubOver", seed = 1234L)

over %>%
  count(default)
# A tibble: 2 x 2
  default      n
  <fct>    <int>
1 No      6767
2 Yes      6767

# over-sampling with random seed, and k = 10
over10 <- sb %>%
  sampling_target(method = "ubOver", seed = 1234L, k = 10)

over10 %>%
  count(default)
# A tibble: 2 x 2
  default      n
  <fct>    <int>
```

```

1 No      6767
2 Yes     2330

# SMOTE with random seed
smote <- sb %>%
  sampling_target(method = "ubSMOTE", seed = 1234L)

smote %>%
  count(default)
# A tibble: 2 x 2
  default      n
  <fct>    <int>
1 No      932
2 Yes     699

# SMOTE with random seed, and perc.under = 250
smote250 <- sb %>%
  sampling_target(method = "ubSMOTE", seed = 1234L, perc.under = 250)

smote250 %>%
  count(default)
# A tibble: 2 x 2
  default      n
  <fct>    <int>
1 No     1165
2 Yes     699

```

The argument that specifies the sampling method in `sampling_target()` is `method`. “ubUnder” is under-sampling, and “ubOver” is over-sampling, “ubSMOTE” is SMOTE (Synthetic Minority Over-sampling Technique).